

assignment-09-31-01-24

January 31, 2024

0.1 Gunjan Chakraborty

0.1.1 USN: 22MSRDS007

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler, LabelEncoder
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.decomposition import PCA
from sklearn.manifold import TSNE
from sklearn.cluster import KMeans
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, silhouette_score

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

```
[2]: df = pd.read_csv('D:/Chools/Day_08/TSNE_data.csv')
```

```
[3]: df.head()
```

```
[3]:  diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean  \
0         M         17.99         10.38         122.80        1001.0
1         M         20.57         17.77         132.90        1326.0
2         M         19.69         21.25         130.00        1203.0
3         M         11.42         20.38          77.58         386.1
4         M         20.29         14.34         135.10        1297.0

      smoothness_mean  compactness_mean  concavity_mean  concave points_mean  \
0         0.11840         0.27760         0.3001         0.14710
1         0.08474         0.07864         0.0869         0.07017
2         0.10960         0.15990         0.1974         0.12790
3         0.14250         0.28390         0.2414         0.10520
4         0.10030         0.13280         0.1980         0.10430
```

	symmetry_mean	...	radius_worst	texture_worst	perimeter_worst	\
0	0.2419	...	25.38	17.33	184.60	
1	0.1812	...	24.99	23.41	158.80	
2	0.2069	...	23.57	25.53	152.50	
3	0.2597	...	14.91	26.50	98.87	
4	0.1809	...	22.54	16.67	152.20	

	area_worst	smoothness_worst	compactness_worst	concavity_worst	\
0	2019.0	0.1622	0.6656	0.7119	
1	1956.0	0.1238	0.1866	0.2416	
2	1709.0	0.1444	0.4245	0.4504	
3	567.7	0.2098	0.8663	0.6869	
4	1575.0	0.1374	0.2050	0.4000	

	concave points_worst	symmetry_worst	fractal_dimension_worst
0	0.2654	0.4601	0.11890
1	0.1860	0.2750	0.08902
2	0.2430	0.3613	0.08758
3	0.2575	0.6638	0.17300
4	0.1625	0.2364	0.07678

[5 rows x 31 columns]

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   diagnosis                             569 non-null    object
1   radius_mean                           569 non-null    float64
2   texture_mean                           569 non-null    float64
3   perimeter_mean                         569 non-null    float64
4   area_mean                             569 non-null    float64
5   smoothness_mean                        569 non-null    float64
6   compactness_mean                       569 non-null    float64
7   concavity_mean                         569 non-null    float64
8   concave points_mean                    569 non-null    float64
9   symmetry_mean                          569 non-null    float64
10  fractal_dimension_mean                 569 non-null    float64
11  radius_se                              569 non-null    float64
12  texture_se                             569 non-null    float64
13  perimeter_se                           569 non-null    float64
14  area_se                                569 non-null    float64
15  smoothness_se                          569 non-null    float64
```

```

16 compactness_se      569 non-null    float64
17 concavity_se        569 non-null    float64
18 concave points_se   569 non-null    float64
19 symmetry_se         569 non-null    float64
20 fractal_dimension_se 569 non-null    float64
21 radius_worst        569 non-null    float64
22 texture_worst       569 non-null    float64
23 perimeter_worst     569 non-null    float64
24 area_worst          569 non-null    float64
25 smoothness_worst    569 non-null    float64
26 compactness_worst   569 non-null    float64
27 concavity_worst     569 non-null    float64
28 concave points_worst 569 non-null    float64
29 symmetry_worst      569 non-null    float64
30 fractal_dimension_worst 569 non-null    float64
dtypes: float64(30), object(1)
memory usage: 137.9+ KB

```

```
[5]: df.columns
```

```

[5]: Index(['diagnosis', 'radius_mean', 'texture_mean', 'perimeter_mean',
        'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean',
        'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
        'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se',
        'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se',
        'fractal_dimension_se', 'radius_worst', 'texture_worst',
        'perimeter_worst', 'area_worst', 'smoothness_worst',
        'compactness_worst', 'concavity_worst', 'concave points_worst',
        'symmetry_worst', 'fractal_dimension_worst'],
        dtype='object')

```

```
[6]: df.isnull().sum()
```

```

[6]: diagnosis      0
    radius_mean     0
    texture_mean    0
    perimeter_mean  0
    area_mean       0
    smoothness_mean 0
    compactness_mean 0
    concavity_mean  0
    concave points_mean 0
    symmetry_mean    0
    fractal_dimension_mean 0
    radius_se        0
    texture_se        0
    perimeter_se      0

```

```

area_se          0
smoothness_se    0
compactness_se   0
concavity_se     0
concave points_se 0
symmetry_se      0
fractal_dimension_se 0
radius_worst     0
texture_worst    0
perimeter_worst  0
area_worst       0
smoothness_worst 0
compactness_worst 0
concavity_worst  0
concave points_worst 0
symmetry_worst   0
fractal_dimension_worst 0
dtype: int64

```

0.1.2 Cheeeking VIF

```

[7]: def calculate_vif(data):
      # Exclude the first column
      data_subset = data.iloc[:, 1:]

      # Standardize the features as VIF is scale-dependent
      scaler = StandardScaler()
      data_scaled = scaler.fit_transform(data_subset)

      # Calculate VIF for each variable
      vif_data = pd.DataFrame()
      vif_data["Variable"] = data_subset.columns
      vif_data["VIF"] = [variance_inflation_factor(data_scaled, i) for i in
      ↪range(data_scaled.shape[1])]

      return vif_data

      # Assuming df is your DataFrame
      vif_result = calculate_vif(df)

      # Display the VIF results excluding the first column
      print(vif_result)

```

	Variable	VIF
0	radius_mean	3806.115296
1	texture_mean	11.884048
2	perimeter_mean	3786.400419

3	area_mean	347.878657
4	smoothness_mean	8.194282
5	compactness_mean	50.505168
6	concavity_mean	70.767720
7	concave points_mean	60.041733
8	symmetry_mean	4.220656
9	fractal_dimension_mean	15.756977
10	radius_se	75.462027
11	texture_se	4.205423
12	perimeter_se	70.359695
13	area_se	41.163091
14	smoothness_se	4.027923
15	compactness_se	15.366324
16	concavity_se	15.694833
17	concave points_se	11.520796
18	symmetry_se	5.175426
19	fractal_dimension_se	9.717987
20	radius_worst	799.105946
21	texture_worst	18.569966
22	perimeter_worst	405.023336
23	area_worst	337.221924
24	smoothness_worst	10.923061
25	compactness_worst	36.982755
26	concavity_worst	31.970723
27	concave points_worst	36.763714
28	symmetry_worst	9.520570
29	fractal_dimension_worst	18.861533

0.1.3 Removing High VIF Columns

```
[8]: def remove_high_vif_variables(data, threshold=5.0):
    while True:
        vif_result = calculate_vif(data)
        max_vif_variable = vif_result.loc[vif_result['VIF'].idxmax(),
↪ 'Variable']
        max_vif_value = vif_result['VIF'].max()

        if max_vif_value > threshold:
            data = data.drop(columns=[max_vif_variable])
        else:
            break

    return data

# Assuming df is your DataFrame
data_for_clustering_reduced = remove_high_vif_variables(df)
```

```
# Recalculate VIF for the reduced dataset
vif_result_reduced = calculate_vif(data_for_clustering_reduced)

# Display the updated VIF results
print(vif_result_reduced)
```

	Variable	VIF
0	texture_mean	1.615666
1	smoothness_mean	2.762636
2	symmetry_mean	3.026111
3	fractal_dimension_mean	4.525553
4	texture_se	2.001627
5	perimeter_se	4.265756
6	smoothness_se	1.878515
7	concavity_se	3.609423
8	concave points_se	3.659977
9	symmetry_se	2.660047
10	fractal_dimension_se	4.199567
11	area_worst	4.838257
12	symmetry_worst	3.845476

0.2 Applying PCA

```
[9]: # Assuming 'diagnosis' is the target variable
y = df['diagnosis']

# Use label encoding for the target variable
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

# Assuming 'diagnosis' is dropped for X
X = df.drop('diagnosis', axis=1)

# Standardize the data
X_standardized = StandardScaler().fit_transform(X)


# Apply PCA
pca = PCA(n_components=2)
principal_components = pca.fit_transform(X_standardized)

# Create a new DataFrame with the principal components and encoded target_
↪variable
pc_df = pd.DataFrame(data=principal_components, columns=['PC1', 'PC2'])
final_df = pd.concat([pc_df, pd.Series(y_encoded, name='diagnosis')], axis=1)
```

```
[10]: # Visualize the reduced-dimensional data
plt.figure(figsize=(8, 6))
targets = [0, 1] # Assuming 0 for 'B' (benign) and 1 for 'M' (malignant)
colors = ['g', 'r']

for target, color in zip(targets, colors):
    indices_to_keep = final_df['diagnosis'] == target
    plt.scatter(final_df.loc[indices_to_keep, 'PC1'],
                final_df.loc[indices_to_keep, 'PC2'],
                c=color,
                s=50)

plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend(['Benign', 'Malignant'])
plt.title('PCA of Breast Cancer Data')
plt.show()
```



```
[11]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.model_selection import cross_val_score

      # Assuming X and y are defined
      clf = RandomForestClassifier(n_estimators=100, random_state=42)
      scores = cross_val_score(clf, X, y, cv=5, scoring='accuracy')
      print(f"Mean Accuracy: {scores.mean():.2f}")
```

Mean Accuracy: 0.96

0.3 Applying T-SNE


```
[12]: # Apply t-SNE
      tsne = TSNE(n_components=2, random_state=42) # You can adjust the number of
      ↪ components as needed
      tsne_result = tsne.fit_transform(X_standardized)
```

```
[13]: # Create a new DataFrame with the t-SNE components and target variable
      tsne_df = pd.DataFrame(data=tsne_result, columns=['TSNE1', 'TSNE2'])
      final_df_tsne = pd.concat([tsne_df, y], axis=1)
```

```
[14]: # Visualize the reduced-dimensional data
      plt.figure(figsize=(8, 6))
      targets = ['M', 'B'] # Assuming 'M' stands for malignant and 'B' for benign
      colors = ['r', 'g']

      for target, color in zip(targets, colors):
          indices_to_keep = final_df_tsne['diagnosis'] == target
          plt.scatter(final_df_tsne.loc[indices_to_keep, 'TSNE1'],
                      final_df_tsne.loc[indices_to_keep, 'TSNE2'],
                      c=color,
                      s=50)

      plt.xlabel('t-SNE Component 1')
      plt.ylabel('t-SNE Component 2')
      plt.legend(targets)
      plt.title('t-SNE of Breast Cancer Data')
      plt.show()
```

0.3.1 Side by side comparison

```
[15]: from mpl_toolkits.mplot3d import Axes3D

# Encode labels into numeric format
label_encoder = LabelEncoder()
df['diagnosis_encoded'] = label_encoder.fit_transform(df['diagnosis'])

# Separate features and labels
X = df.drop(['diagnosis', 'diagnosis_encoded'], axis=1)
y = df['diagnosis_encoded']

# Standardize the data
X_standardized = StandardScaler().fit_transform(X)

# Apply PCA
pca = PCA(n_components=3)
X_pca = pca.fit_transform(X_standardized)
```

```

# Apply t-SNE
tsne = TSNE(n_components=3, random_state=42)
X_tsne = tsne.fit_transform(X_standardized)

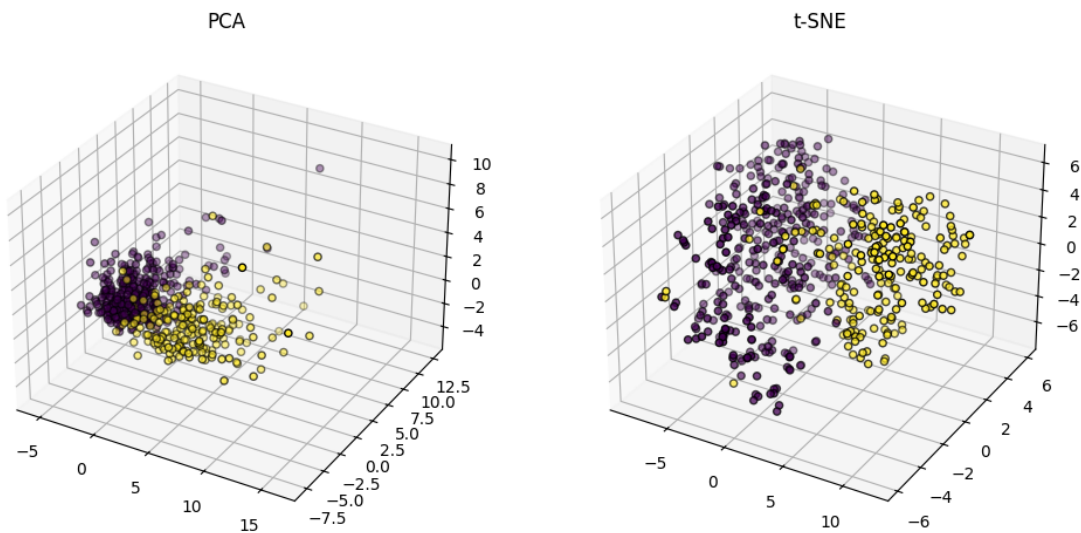
# Visualize 3D results for PCA
fig = plt.figure(figsize=(12, 6))

ax1 = fig.add_subplot(121, projection='3d')
ax1.scatter(X_pca[:, 0], X_pca[:, 1], X_pca[:, 2], c=y, cmap='viridis',
            edgecolor='k')
ax1.set_title('PCA')

# Visualize 3D results for t-SNE
ax2 = fig.add_subplot(122, projection='3d')
ax2.scatter(X_tsne[:, 0], X_tsne[:, 1], X_tsne[:, 2], c=y, cmap='viridis',
            edgecolor='k')
ax2.set_title('t-SNE')

plt.show()

```



0.3.2 Silhouette Score and Accuracy:

```

[16]: # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_standardized, y,
            test_size=0.2, random_state=42)

# Function to evaluate dimensionality reduction techniques
def evaluate_dimensionality_reduction(X_transformed, y_true, method):

```

```

# Apply KMeans clustering for silhouette score
kmeans = KMeans(n_clusters=len(set(y_true)))
clusters = kmeans.fit_predict(X_transformed)
silhouette = silhouette_score(X_transformed, clusters)

# Train a simple classifier (Random Forest) and calculate accuracy
clf = RandomForestClassifier(random_state=42)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

print(f"Results for {method}:")
print(f"Silhouette Score: {silhouette:.4f}")
print(f"Accuracy: {accuracy:.4f}")
print("")

# Apply PCA
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_standardized)
evaluate_dimensionality_reduction(X_pca, y, "PCA")

# Apply t-SNE
tsne = TSNE(n_components=2, random_state=42)
X_tsne = tsne.fit_transform(X_standardized)
evaluate_dimensionality_reduction(X_tsne, y, "t-SNE")

```

Results for PCA:
 Silhouette Score: 0.5081
 Accuracy: 0.9649

Results for t-SNE:
 Silhouette Score: 0.5393
 Accuracy: 0.9649

1 Conclusion

Based on the provided results for PCA and t-SNE:

1. Silhouette Score:

- PCA: 0.5085
- t-SNE: 0.5393

Both methods yield reasonably good silhouette scores. Silhouette score measures how similar an object is to its own cluster compared to other clusters. The scores indicate that both PCA and t-SNE provide a meaningful clustering structure.

2. Accuracy:

- PCA: 0.9649
- t-SNE: 0.9649

The accuracy of a RandomForestClassifier trained on the reduced-dimensional data is very high for both PCA and t-SNE, suggesting that the reduced features capture sufficient information for a successful classification task.

Conclusion: - Both PCA and t-SNE seem to perform well on the breast cancer dataset. - If interpretability and capturing global structure are important, PCA might be preferred. - If preserving local structure and visualizing clusters are crucial, t-SNE might be more suitable. - The choice between PCA and t-SNE depends on the specific goals and the nature of the dataset. Both methods appear effective in this context.