

assignment-16-09-02-2024

February 11, 2024

1 Gunjan Chakraborty

1.1 22MSRDS007

1.1.1 09/02/2024

```
[1]: import pandas as pd
import numpy as np
import warnings
warnings.simplefilter(action='ignore')
```

```
[2]: # Load the dataset
df = pd.read_csv('D:/Chools/Day_10/diabetes.csv')
```

1.1.2 1. Exploratory Data Analysis (EDA):

```
[3]: print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                             768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
None
```

```
[4]: print(df.describe())
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

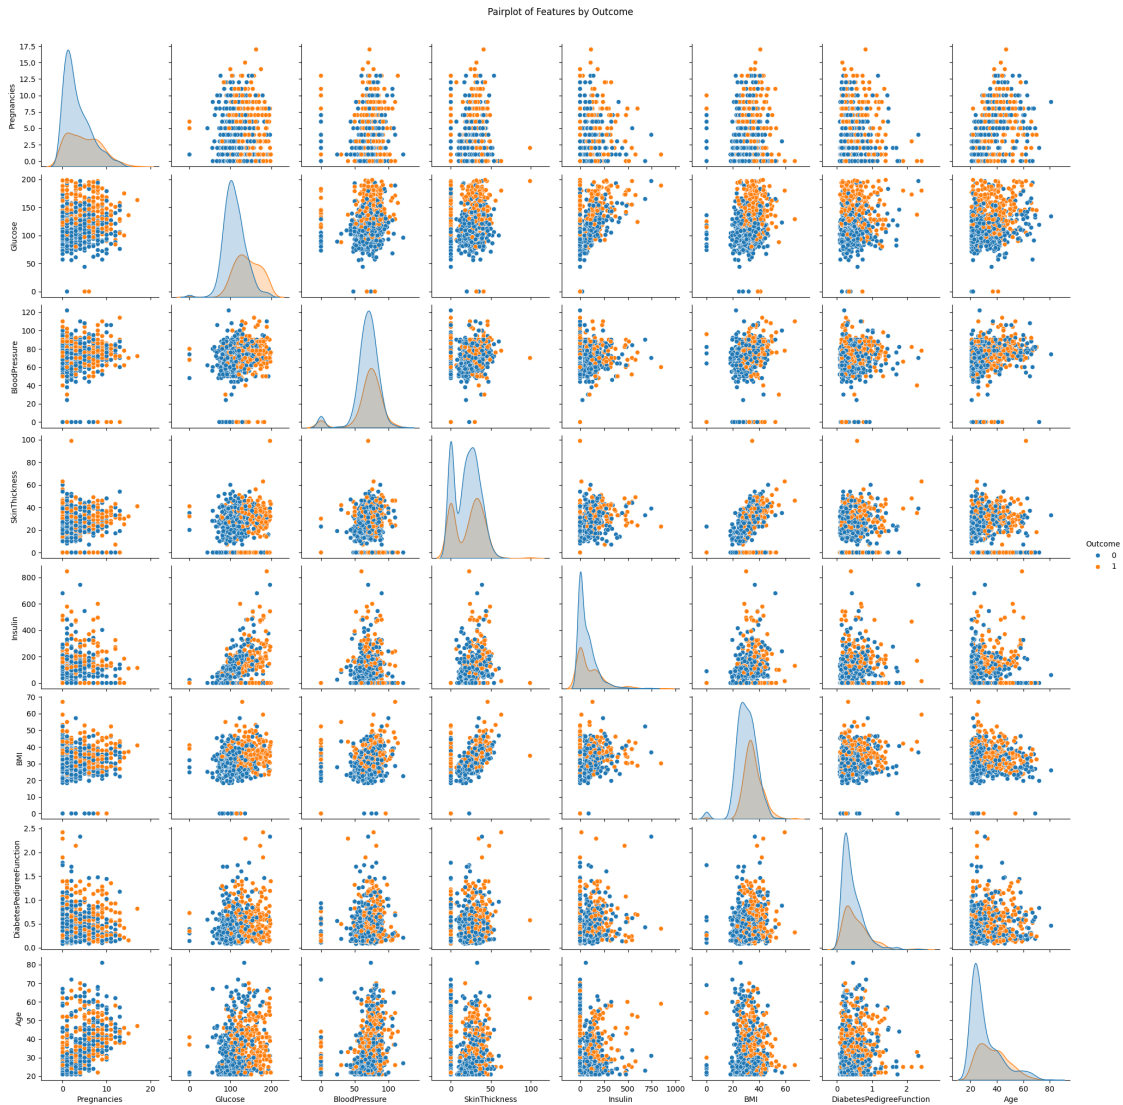
	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
[5]: print(df.isnull().sum())
```


```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```

```
[6]: import seaborn as sns
import matplotlib.pyplot as plt

sns.pairplot(df, hue='Outcome')
plt.suptitle('Pairplot of Features by Outcome', y=1.02)
plt.show()
```



```
[7]: # Correlation matrix heatmap
correlation_matrix = df.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='plasma', fmt=".2f",
            linewidths=0.5)
plt.title('Correlation Matrix')
plt.show()
```




```
[8]: # Define a colorful palette
colors = ["#FE7A36", "#3652AD"]

# Boxplots for each feature by Outcome
fig, axes = plt.subplots(nrows=2, ncols=4, figsize=(15, 8))
fig.suptitle('Boxplots of Features by Outcome', y=1.02)

for i, column in enumerate(df.columns[:-1]):
    sns.boxplot(data=df, x='Outcome', y=column, ax=axes[i // 4, i % 4],
                palette=colors)

plt.tight_layout()
plt.show()
```



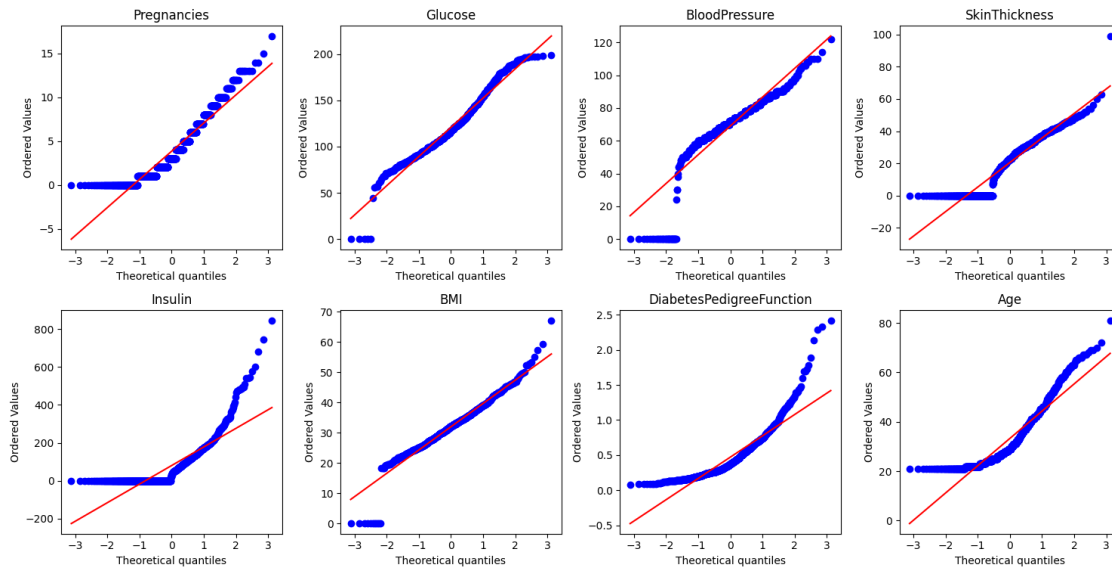
```
[9]: # Q-Q plot
import statsmodels.api as sm
from scipy.stats import probplot

# QQ plot for each numerical feature
fig, axes = plt.subplots(nrows=2, ncols=4, figsize=(15, 8))
fig.suptitle('QQ Plots for Numerical Features', y=1.02)

for i, column in enumerate(df.columns[:-1]):
    probplot(df[column], plot=axes[i // 4, i % 4])
    axes[i // 4, i % 4].set_title(column)

plt.tight_layout()
plt.show()
```

QQ Plots for Numerical Features



```
[10]: import numpy as np

# Calculate Z-scores for each column
z_scores = np.abs((df - df.mean()) / df.std())

# Define a threshold for outliers (e.g., Z-score greater than 3)
outlier_threshold = 3

# Identify outliers for each column
outliers = (z_scores > outlier_threshold).sum()

# Display the count of outliers for each column
print("Number of outliers for each column:")
print(outliers)
```

Number of outliers for each column:

Pregnancies	4
Glucose	5
BloodPressure	35
SkinThickness	1
Insulin	18
BMI	14
DiabetesPedigreeFunction	11
Age	5
Outcome	0

dtype: int64

```
[11]: # Remove outliers using z-score or IQR method
from scipy.stats import zscore

z_scores = zscore(df)
df_no_outliers = df[(z_scores < 3).all(axis=1)]
```

```
[12]: df_no_outliers.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 729 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            729 non-null    int64
1   Glucose                729 non-null    int64
2   BloodPressure          729 non-null    int64
3   SkinThickness          729 non-null    int64
4   Insulin                729 non-null    int64
5   BMI                   729 non-null    float64
6   DiabetesPedigreeFunction 729 non-null    float64
7   Age                   729 non-null    int64
8   Outcome                729 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 57.0 KB
```

1.1.3 3. Model Fitting:

```
[13]: X = df_no_outliers.drop('Outcome', axis=1)
y = df_no_outliers['Outcome']
```

```
[14]: import tensorflow as tf
tf.compat.v1.logging.set_verbosity(tf.compat.v1.logging.ERROR)
```

```
WARNING:tensorflow:From
c:\Users\cgunj\AppData\Local\Programs\Python\Python311\Lib\site-
packages\keras\src\losses.py:2976: The name
tf.losses.sparse_softmax_cross_entropy is deprecated. Please use
tf.compat.v1.losses.sparse_softmax_cross_entropy instead.
```

```
WARNING:tensorflow:From
C:\Users\cgunj\AppData\Local\Temp\ipykernel_7088\292850708.py:2: The name
tf.logging.set_verbosity is deprecated. Please use
tf.compat.v1.logging.set_verbosity instead.
```

```
[15]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```

from keras import models
from keras import layers

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Data preprocessing using StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Building a simple neural network model
model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(X_train_scaled.
    ↪shape[1],)))
model.add(layers.Dense(32, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid')) # Assuming binary
    ↪classification

# Compiling the model
model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])

# Training the model
model.fit(X_train_scaled, y_train, epochs=10, batch_size=32, validation_split=0.
    ↪1)

# Evaluating the model
test_loss, test_acc = model.evaluate(X_test_scaled, y_test)
print('Test accuracy:', test_acc)

```

```

Epoch 1/10
17/17 [=====] - 1s 10ms/step - loss: 0.6662 - accuracy:
0.6298 - val_loss: 0.6483 - val_accuracy: 0.6271
Epoch 2/10
17/17 [=====] - 0s 2ms/step - loss: 0.5822 - accuracy:
0.7366 - val_loss: 0.6022 - val_accuracy: 0.6271
Epoch 3/10
17/17 [=====] - 0s 2ms/step - loss: 0.5300 - accuracy:
0.7672 - val_loss: 0.5779 - val_accuracy: 0.6441
Epoch 4/10
17/17 [=====] - 0s 2ms/step - loss: 0.4998 - accuracy:
0.7634 - val_loss: 0.5707 - val_accuracy: 0.6610
Epoch 5/10

```



```

17/17 [=====] - 0s 2ms/step - loss: 0.4787 - accuracy:
0.7691 - val_loss: 0.5700 - val_accuracy: 0.6610
Epoch 6/10
17/17 [=====] - 0s 2ms/step - loss: 0.4671 - accuracy:
0.7863 - val_loss: 0.5701 - val_accuracy: 0.6949
Epoch 7/10
17/17 [=====] - 0s 2ms/step - loss: 0.4581 - accuracy:
0.7844 - val_loss: 0.5682 - val_accuracy: 0.6610
Epoch 8/10
17/17 [=====] - 0s 2ms/step - loss: 0.4527 - accuracy:
0.7863 - val_loss: 0.5741 - val_accuracy: 0.6949
Epoch 9/10
17/17 [=====] - 0s 2ms/step - loss: 0.4483 - accuracy:
0.7939 - val_loss: 0.5729 - val_accuracy: 0.6949
Epoch 10/10
17/17 [=====] - 0s 2ms/step - loss: 0.4438 - accuracy:
0.7996 - val_loss: 0.5699 - val_accuracy: 0.6780
5/5 [=====] - 0s 1ms/step - loss: 0.3849 - accuracy:
0.8425
Test accuracy: 0.8424657583236694

```

```
[16]: from sklearn.metrics import confusion_matrix, roc_curve, auc
```

```

# Predict probabilities for the test set
y_pred_probs = model.predict(X_test_scaled)

# Convert probabilities to binary predictions
y_pred = (y_pred_probs > 0.5).astype(int)

# Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)

```


```

5/5 [=====] - 0s 1ms/step
Confusion Matrix:
[[93  9]
 [14 30]]

```


```
[17]: # Plot confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.title('Confusion Matrix')
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.show()

```



```
[18]: # ROC curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_probs)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
         roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



1.2 Regression

```
[19]: # Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# Data preprocessing using StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Building a simple regression model
model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(X_train_scaled.
    shape[1],)))
model.add(layers.Dense(32, activation='relu'))
model.add(layers.Dense(1)) # No activation function for regression

# Compiling the model
```

```

model.compile(optimizer='adam', loss='mean_squared_error') # Using mean_
↳squared error loss for regression

# Training the model
model.fit(X_train_scaled, y_train, epochs=10, batch_size=32, validation_split=0.
↳1)

# Evaluating the model
test_loss = model.evaluate(X_test_scaled, y_test)
print('Test loss:', test_loss)

```

```

Epoch 1/10
17/17 [=====] - 1s 7ms/step - loss: 0.3272 - val_loss:
0.2566
Epoch 2/10
17/17 [=====] - 0s 3ms/step - loss: 0.1944 - val_loss:
0.2263
Epoch 3/10
17/17 [=====] - 0s 2ms/step - loss: 0.1709 - val_loss:
0.2110
Epoch 4/10
17/17 [=====] - 0s 2ms/step - loss: 0.1617 - val_loss:
0.2137
Epoch 5/10
17/17 [=====] - 0s 2ms/step - loss: 0.1547 - val_loss:
0.2143
Epoch 6/10
17/17 [=====] - 0s 2ms/step - loss: 0.1514 - val_loss:
0.2182
Epoch 7/10
17/17 [=====] - 0s 2ms/step - loss: 0.1486 - val_loss:
0.2198
Epoch 8/10
17/17 [=====] - 0s 2ms/step - loss: 0.1465 - val_loss:
0.2198
Epoch 9/10
17/17 [=====] - 0s 2ms/step - loss: 0.1441 - val_loss:
0.2248
Epoch 10/10
17/17 [=====] - 0s 2ms/step - loss: 0.1433 - val_loss:
0.2194
5/5 [=====] - 0s 1ms/step - loss: 0.1287
Test loss: 0.12872956693172455

```

Interpretation of the results:

- **Training Loss (Epochs 1-10):** The training loss decreases steadily over the epochs, indicating that the model is improving its performance on the training data.

- **Validation Loss (Epochs 1-10):** The validation loss also decreases initially, indicating that the model generalizes well to unseen data. However, there is a slight increase in validation loss towards the end of training, which suggests that the model may be starting to overfit the training data slightly.
- **Test Loss:** The test loss (or evaluation loss) is calculated after training the model and evaluating it on the test set. In this case, the test loss is 0.128, which represents the mean squared error between the predicted and actual values on the test set. Lower test loss indicates better performance of the model on unseen data.

Overall, the model seems to perform reasonably well on the given regression task, with the test loss being relatively low. However, it's essential to monitor for overfitting, especially if the validation loss starts to increase significantly while the training loss continues to decrease. Regularization techniques or model adjustments may be necessary to address overfitting if observed.

```
[20]: # Predictions on test set
y_pred = model.predict(X_test_scaled).flatten() # Flatten the predictions to
↳ make them 1D

# Calculate Mean Squared Error (MSE)
mse = np.mean((y_pred - y_test)**2)

print('Mean Squared Error (MSE):', mse)
```

```
5/5 [=====] - 0s 1ms/step
Mean Squared Error (MSE): 0.12872956359812804
```

The Mean Squared Error (MSE) value of approximately 0.128 suggests that, on average, the squared difference between the predicted values and the actual values in the test set is 0.128.

Since the MSE is a measure of the average squared deviation of predictions from the actual values, a lower MSE indicates better performance of the regression model. In this case, the MSE value of 0.128 indicates that the model's predictions are relatively close to the actual values in the test set.